

Medical AI and deep learning for precision imaging

a hands-on example

PRESIMAL 2022
AUTUMN RESEARCH SCHOOL
AI IN MEDICAL IMAGING
15.09.22



Western Norway
University of
Applied Sciences



Haukeland University Hospital

Alexander Selvikvåg Lundervold

allu@hvl.no; lundervold.net

Arvid Lundervold

arvid.Lundervold@uib.no

“myScript.sh”

Hauke Bartsch

PhD. Ass.prof at
University of
Bergen. Senior
Researcher at
HUS.





Search or jump to...

/ Pull requests Issues Marketplace Explore



MMIV-ML / presimal2022

Public

Edit Pins

Unwatch 2

Fork 1

Star 3

<> Code Issues Pull requests Actions Projects Wiki Security Insights Settings

master

2 branches

0 tags

Go to file

Add file

Code

alu042 upd

13c7b9d 5 hours ago 21 commits

assets

upd

5 hours ago

notebooks

upd

5 hours ago

.gitignore

ignores

2 days ago

LICENSE

ignore & license added

10 days ago

README.md

Update README.md

23 hours ago

environment.yml

conda env

7 days ago

README.md



PRESIMAL 2022 workshop



About



Material from a workshop in the
PRESIMAL Autumn Research School
2022

Readme

GPL-3.0 license

3 stars

2 watching

1 fork

Releases

No releases published

Create a new release

Packages

No packages published

Publish your first package

Languages

Jupyter Notebook 100.0%

Plan



Hands-on



What is deep learning?



*How do you *do* deep learning?*



Some perspectives,
opportunities and
challenges

Deep learning

searching for good hierarchical geometric representations

Deep learning

searching for good hierarchical geometric representations

Deep learning

*searching for **good** hierarchical geometric representations*

Deep learning

*searching for good **hierarchical geometric representations***

Deep learning

searching for good hierarchical geometric representations

01

02

03

04

05

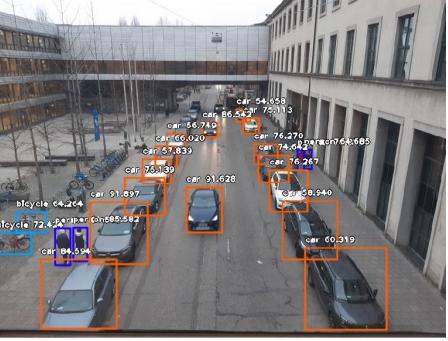
Function approximation

Function approximation

$$y \approx f(x; \theta)$$

what's in the image

an image



Function approximation

$$y \approx f(x; \theta)$$

what's in the image

an image

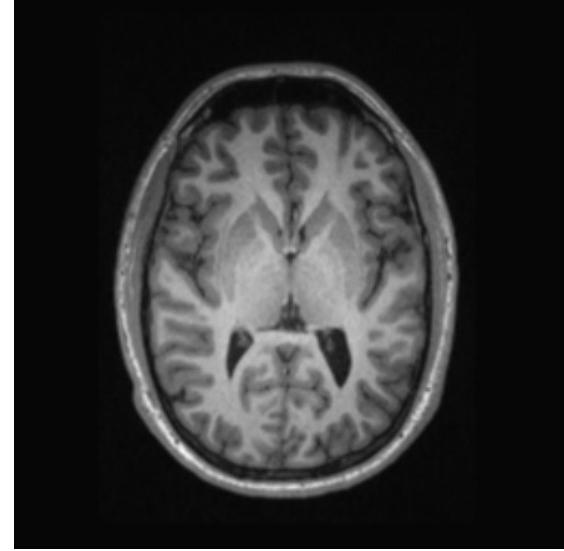
Function approximation

$$y \approx f(x; \theta)$$

what's in the image

an image

AGE

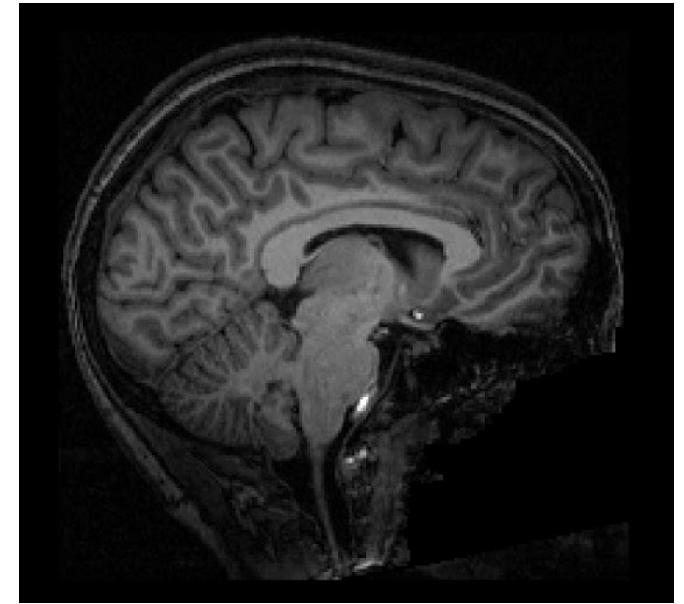
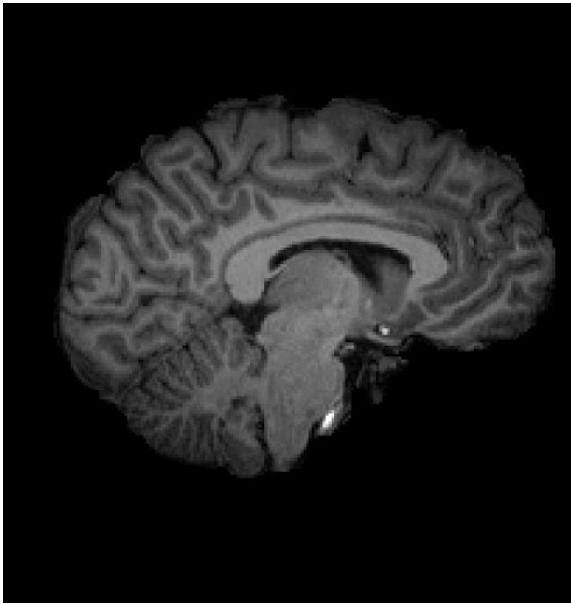


Function approximation

$$y \approx f(x; \theta)$$

"Brain age"

an MRI image



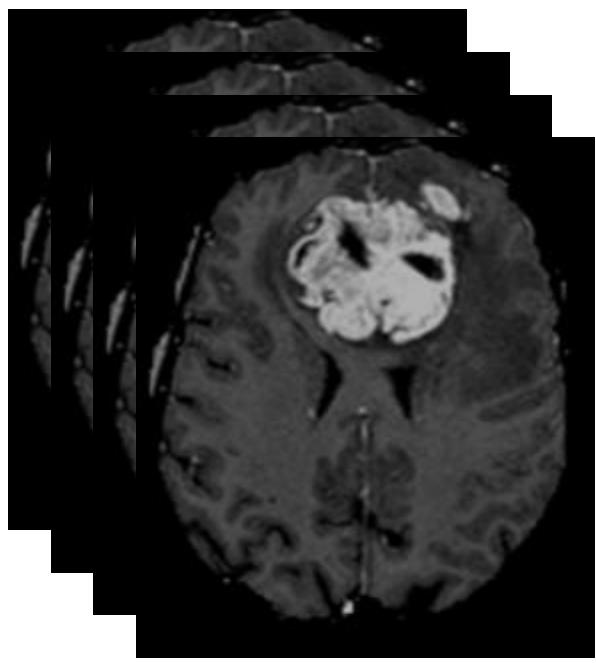
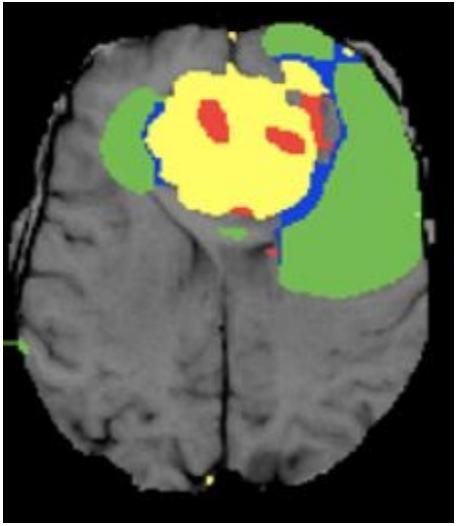
Function approximation

$$y \approx f(x; \theta)$$

skull-stripped MRI

a brain MRI

Diagram illustrating function approximation. A mathematical equation $y \approx f(x; \theta)$ is shown, where x is labeled as "skull-stripped MRI" and y is labeled as "a brain MRI". Arrows point from the labels to the corresponding terms in the equation.

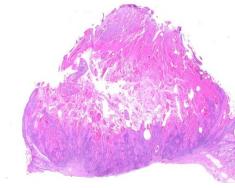
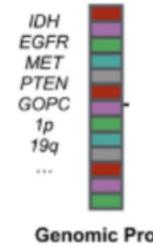
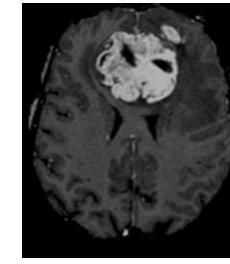
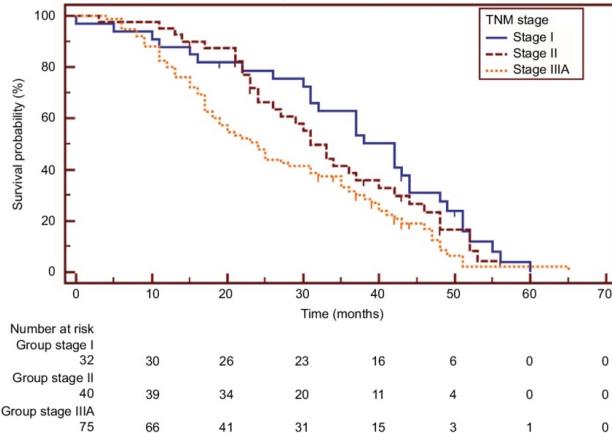


Function approximation

Segmentation of region-of-interests (e.g., tumors)

$$y \approx f(x; \theta)$$

multimodal MRI images



	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q
1.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
2.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
3.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
4.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
5.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
6.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
7.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
8.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
9.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
10.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
11.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
12.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
13.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
14.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
15.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
16.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
17.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
18.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
19.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%
20.	11/10/2000	Adenocarcinoma	3	3	3	3	3	3	338	42	36.8	172	8	0	15%	8%	82%

Function approximation

$$y \approx f(x; \theta)$$

outcome (e.g., survival)

heterogeneous information

Training

$$(X, y)$$

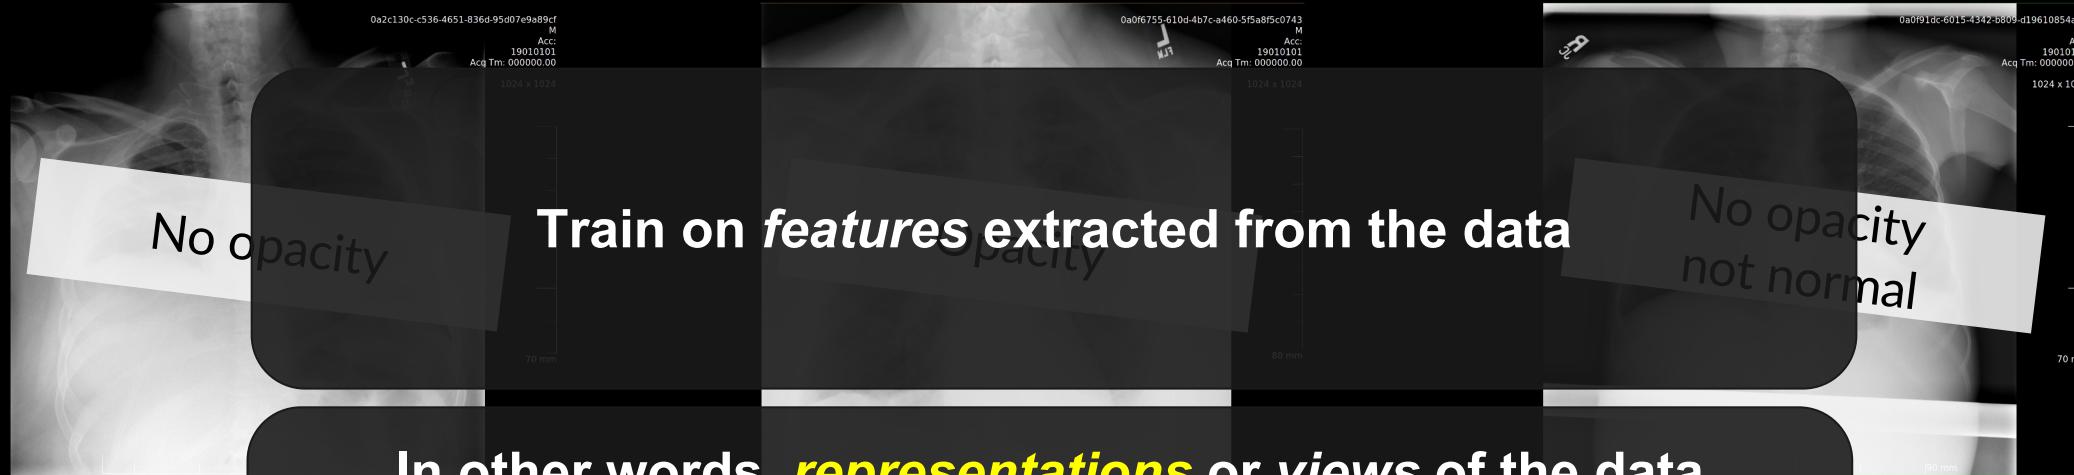
Training

$$(X, y)$$

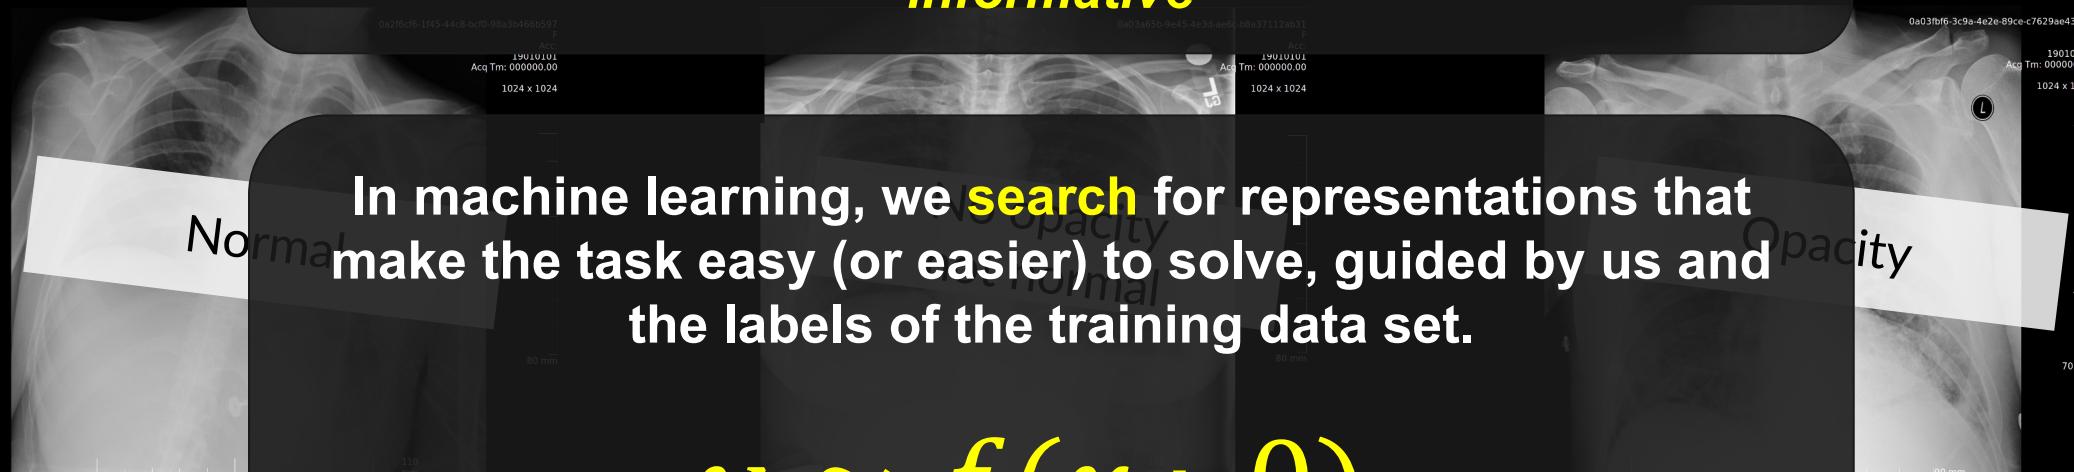
Input data

Labels

*This is the setup in what's called **supervised learning***



Important that these representations are sufficiently informative



$$y \approx f(x ; \theta)$$

01 Machine learning: **function approximation** based on **training**

$$y \approx f(x ; \theta)$$

features of a disease,
process, patient or
system

healthy or not healthy /
outcome

The diagram shows a mathematical equation $y \approx f(x ; \theta)$ centered on a dark background. Two yellow arrows point from text labels to specific parts of the equation. One arrow points from the label 'features of a disease, process, patient or system' to the term x . Another arrow points from the label 'healthy or not healthy / outcome' to the term y .

Models and parameters

$$y \approx f(x; \theta)$$


The diagram illustrates the components of a machine learning model. The equation $y \approx f(x; \theta)$ is centered. Two arrows point from the words "model" and "parameters" at the bottom to the function f and the parameter θ respectively.

***Training* is a search for useful representations through a space defined by a family \mathcal{F} of parametrized models**

01

Machine learning: **function approximation** based on **training**

02

We pick the model family \mathcal{F} . The parameters Θ are found automatically through training

$$y \approx f(x; \theta)$$

Theory versus practice

Theory $X \longrightarrow Y$

In theory:

can get arbitrarily good approximations by training simple models
(*universal approximation and no free lunch theorems*)

Practice $X \xrightarrow{\text{complex, wavy function}} Y$

In practice:

which f we use and how the training is done makes a huge difference

The family \mathcal{F} of models considered, i.e., the *hypothesis space*, should be
(i) **efficiently searchable**, (ii) leading to **expressive** models that can
(iii) **generalize** beyond the training data distribution (i.e., work well on new data)

01

Machine learning: **function approximation** based on **training**

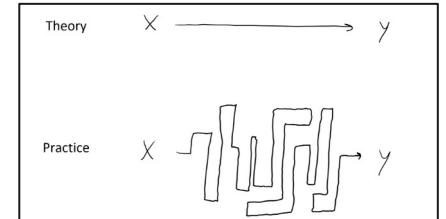
02

We pick the model family \mathcal{F} . The parameters Θ are found automatically through training

03

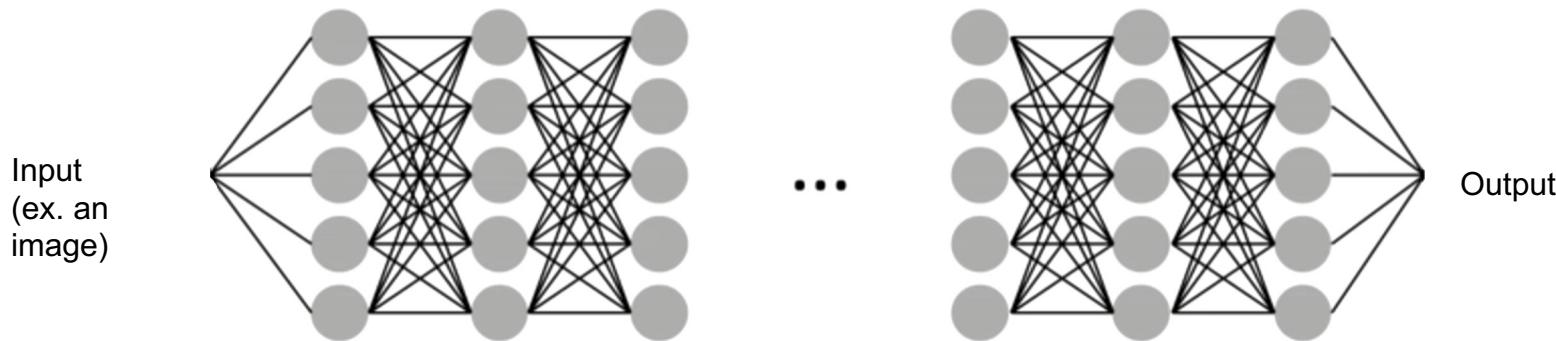
In practice: some models are better suited to some tasks than others.

$$y \approx f(x; \theta)$$



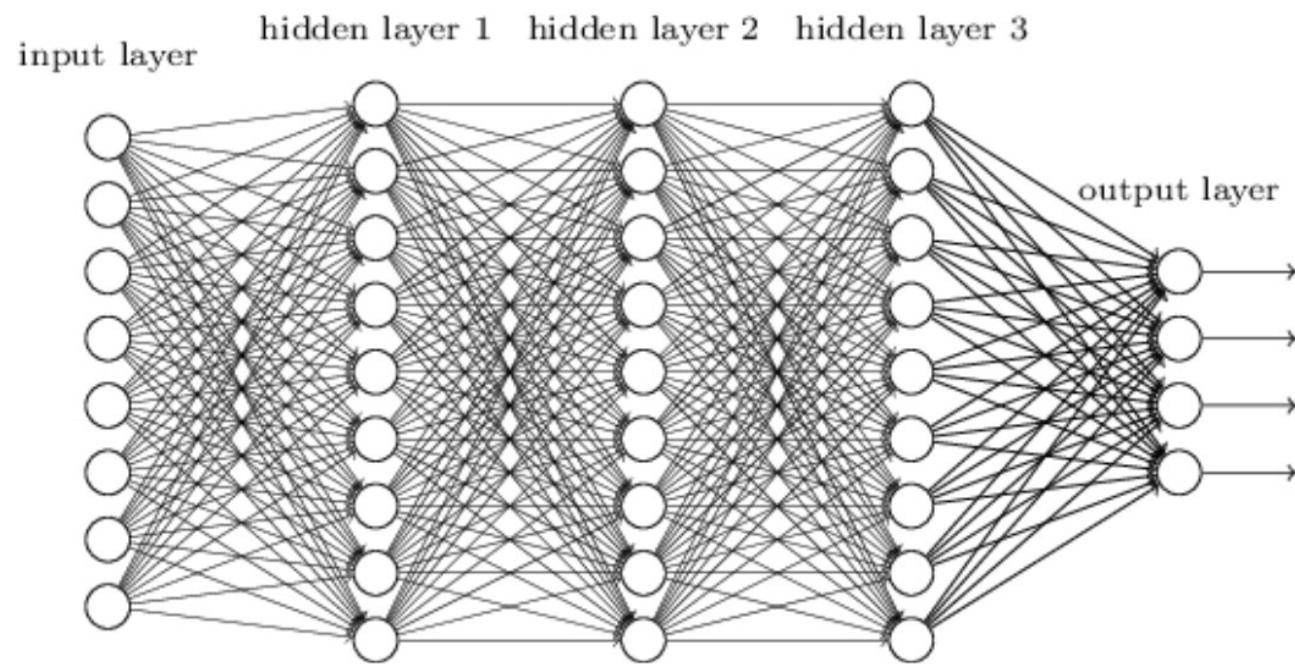
Deep learning

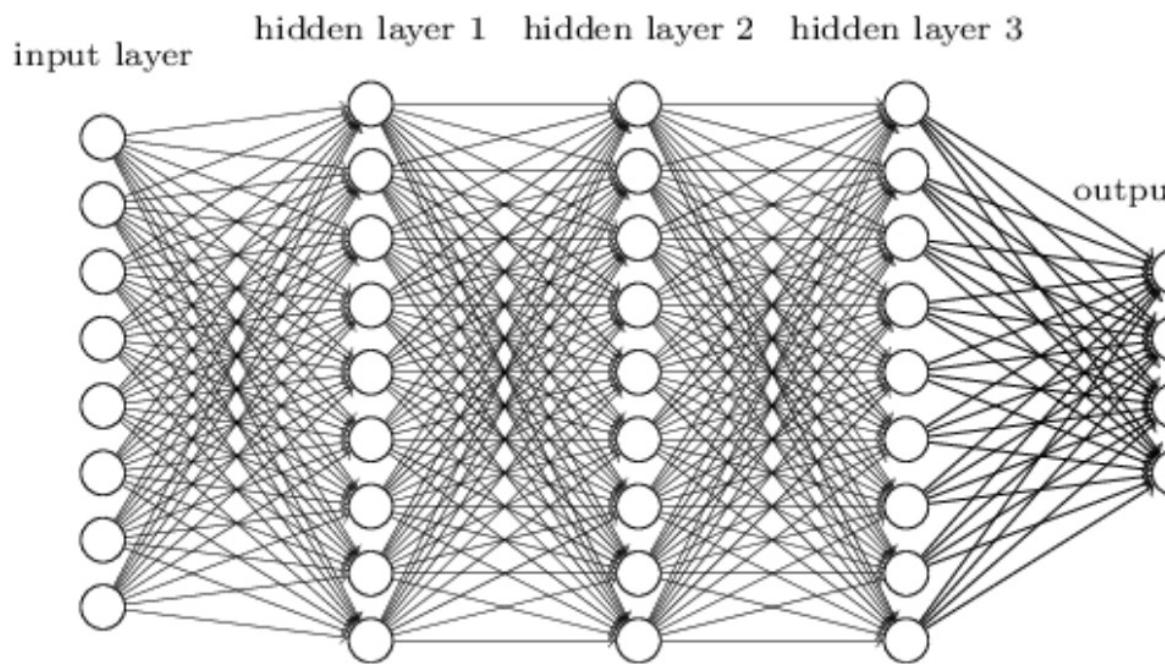
$$y \approx f(x; \theta)$$



Artificial neural networks

Computational graphs of simple units (“neurons”) connected in various specific ways, parametrized by parameters associated to each layer.





A geometric analogy:
uncrumpling paper



This is a fancy way to draw what's really just a composition of simple functions:

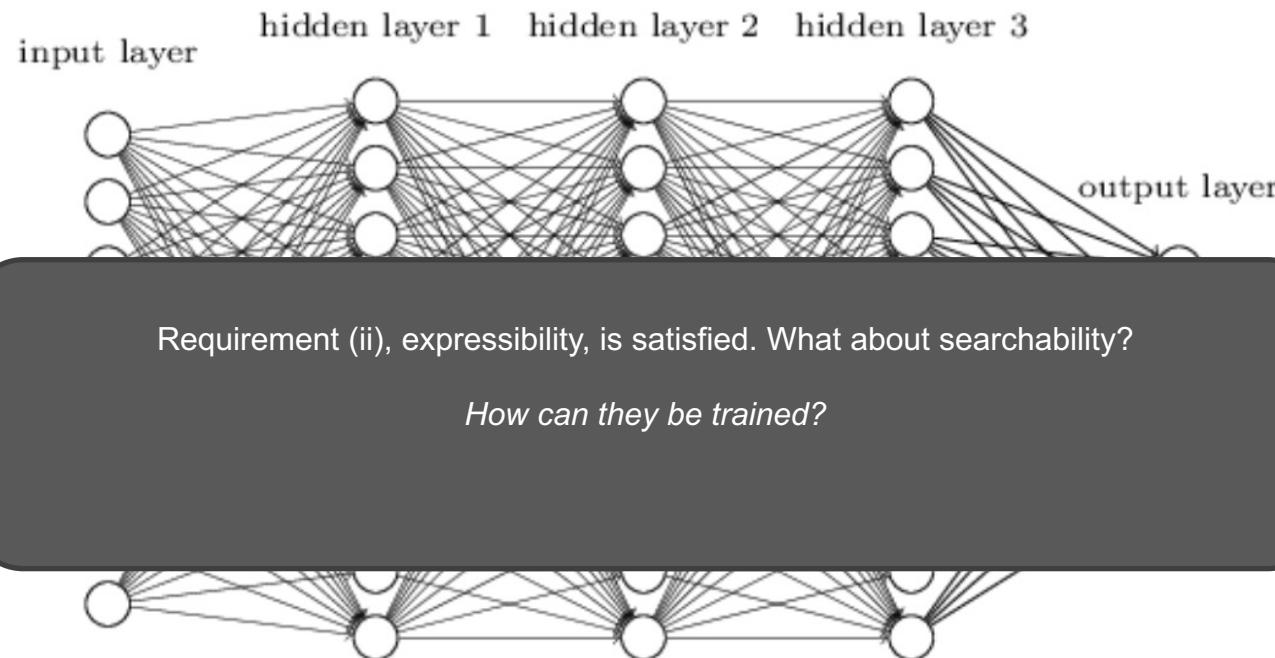
$$y = \sigma_{k+1}(W_{k+1}^T \sigma_k(W_k^T \left(\dots \left(\sigma_1(W_1^T x) \right) \dots \right)) = f(x; \theta)$$

Input x is sent through several consecutive **layers** (functions). An output is produced.

Think of each layer as a (geometric) **transformation** of its input.

Every layer has a bunch of **parameters**, stored in the matrices W

Each family member of F is complex transformation broken down into a composition of simpler ones, parametrized by the layer weights.



This is a fancy way to draw what's really just a composition of simple functions:

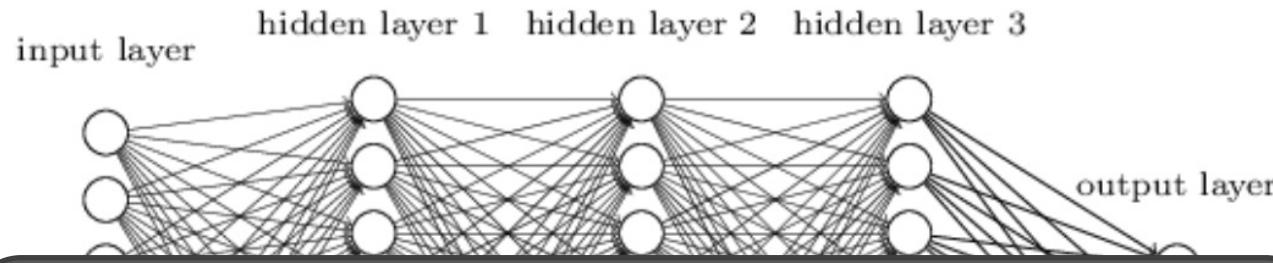
$$y = \sigma_{k+1}(W_{k+1}^T \sigma_k(W_k^T \left(\dots \left(\sigma_1(W_1^T x) \right) \dots \right)) = f(x; \theta)$$

Input x is sent through several consecutive **layers** (functions). An output is produced.

Think of each layer as a (geometric) **transformation** of its input.

Every layer has a bunch of **parameters**, stored in the matrices W

Each family member of \mathcal{F} is complex transformation broken down into a composition of simpler ones, parametrized by the layer weights.



Requirement (ii), expressibility, is satisfied. What about searchability?

How can they be trained?

Concretely, how are the parameters associated to each neuronal connection updated?



This is a fancy way to draw what's really just a composition of simple functions:

$$y = \sigma_{k+1}(W_{k+1}^T \sigma_k(W_k^T \left(\dots \left(\sigma_1(W_1^T x) \right) \dots \right)) = f(x; \theta)$$

Input x is sent through several consecutive **layers** (functions). An output is produced.

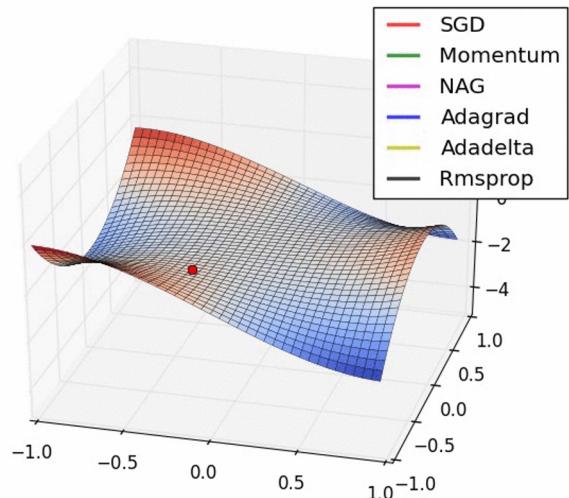
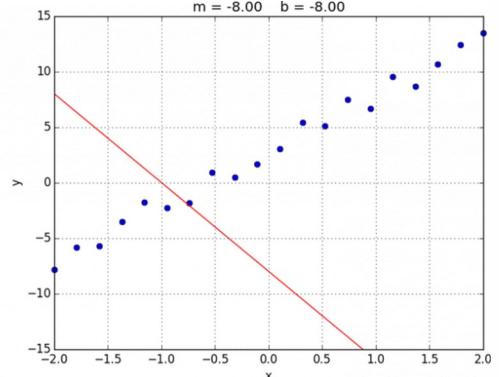
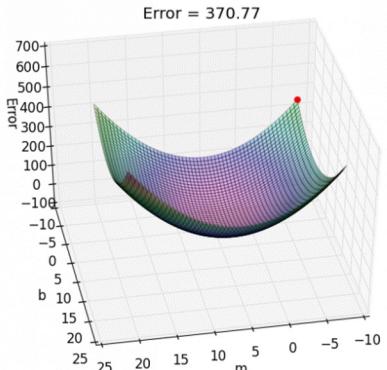
Think of each layer as a (geometric) **transformation** of its input.

Every layer has a bunch of **parameters**, stored in the matrices W

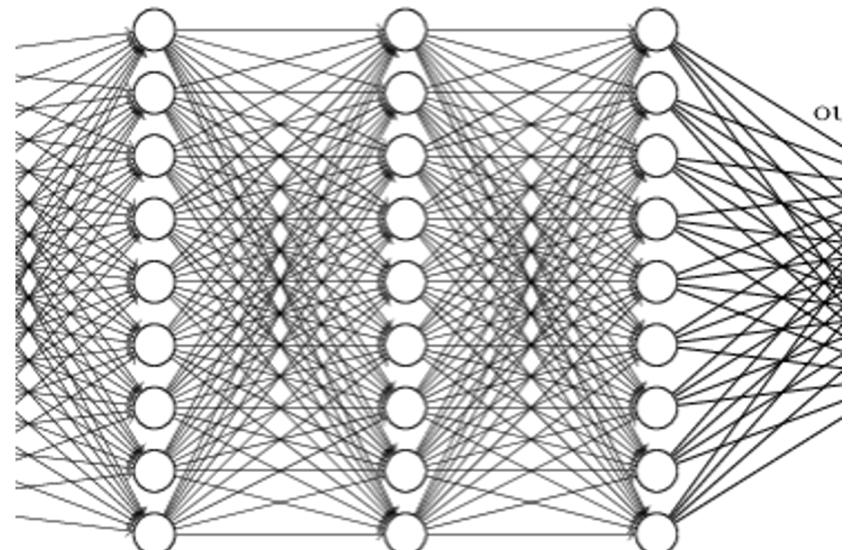
Each family member of \mathcal{F} is complex transformation broken down into a composition of simpler ones, parametrized by the layer weights.

Three basic ingredients:

1. A **loss function**: Used to measure the network's performance on the training data. A **feedback signal** for the training process.
2. An **optimizer**: Used to update the *parameters* of the network to increase performance as measured by the loss function.
Gradient descent and backpropagation.
3. One or more **metrics**: A way to score the model. For example *accuracy* or *mean squared error*.



hidden layer 1 hidden layer 2 hidden layer 3



Define a neural network model with an architecture and a set of trainable parameters (weights and biases)

1. Initialize the parameters

2. Collect a batch of inputs (X, y) and use the model with its current parameters to make predictions y_{pred}

3. Calculate how good the model is using the loss function to measure the discrepancy between the known labels y_{true} and the predictions y_{pred} for this batch

4. Calculate the gradients for each parameter using backpropagation. This measures how changing each parameter would change the loss. In other words, each parameter's contribution to the loss.

5. Move each weight a little bit in the opposite direction of its gradient using gradient descent. I.e., some variant of
$$\text{parameter} = \text{parameter} - \text{learning_rate} * \text{gradient}$$

6. Go back to step 2 and repeat the process

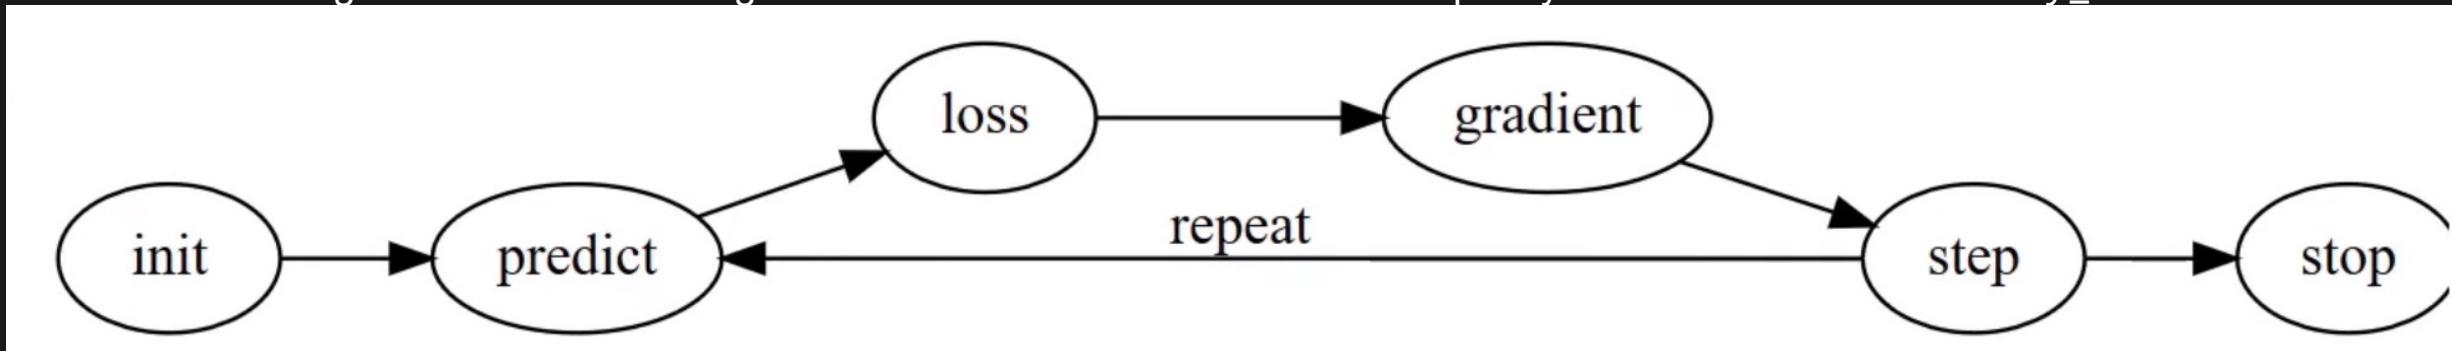
7. Do this over and over for several epochs (one epoch is one pass through all the training data), until you meet some chosen stopping criteria

Define a neural network model with an architecture and a set of trainable parameters (weights and biases)

1. Initialize the parameters

2. Collect a batch of inputs (X, y) and use the model with its current parameters to make predictions y_{pred}

3. Calculate how good the model is using the loss function to measure the discrepancy between the known labels y_{true} and the



5. Move each weight a little bit in the opposite direction of its gradient using gradient descent. i.e., some variant of
$$\text{parameter} = \text{parameter} - \text{learning_rate} * \text{gradient}$$

6. Go back to step 2 and repeat the process

7. Do this over and over for several epochs (one epoch is one pass through all the training data), until you meet some chosen stopping criteria

Define a neural network model with an architecture and a set of trainable parameters (weights and biases)

1. Initialize the parameters
2. Collect a batch of inputs (X, y) and use the model with its current parameters to make predictions y_{pred}
3. Calculate how good the model is using the loss function to measure the discrepancy between the known labels y_{true} and the predictions y_{pred} for this batch
4. Calculate the gradients for each parameter using backpropagation. This measures how changing each parameter would change the loss. In other words, each parameter's contribution to the loss.
5. Move each weight a little bit in the opposite direction of its gradient using gradient descent. I.e., some variant of
$$\text{parameter} = \text{parameter} - \text{learning_rate} * \text{gradient}$$
6. Go back to step 2 and repeat the process
7. Do this over and over for several epochs (one epoch is one pass through all the training data), until you meet some chosen stopping criteria

What is gradient descent?

Define a neural network model with an architecture and a set of trainable parameters (weights and biases)

1. Initialize

2. Collect

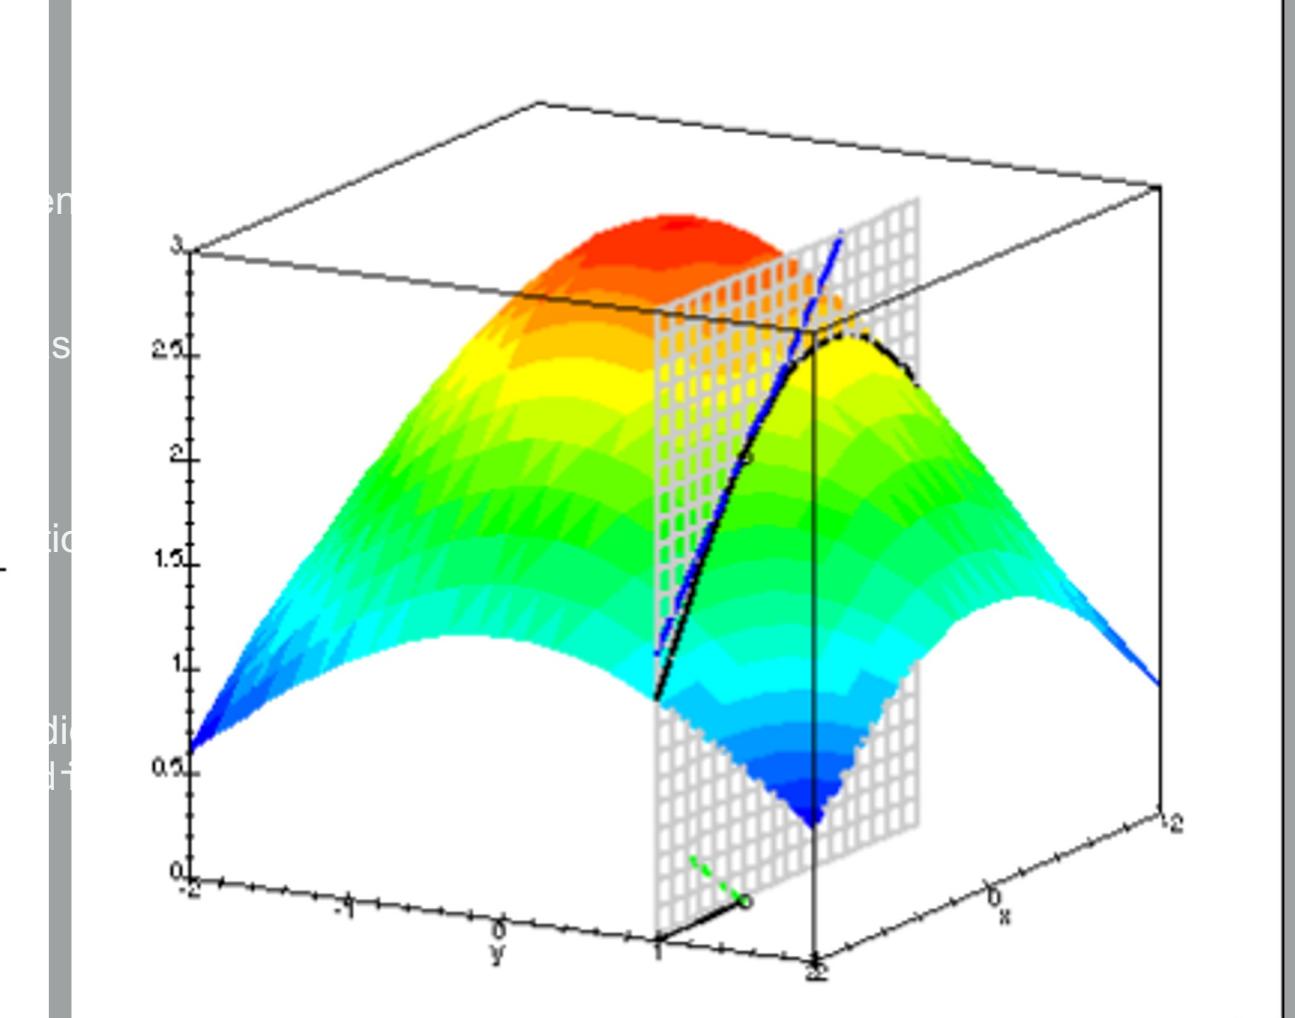
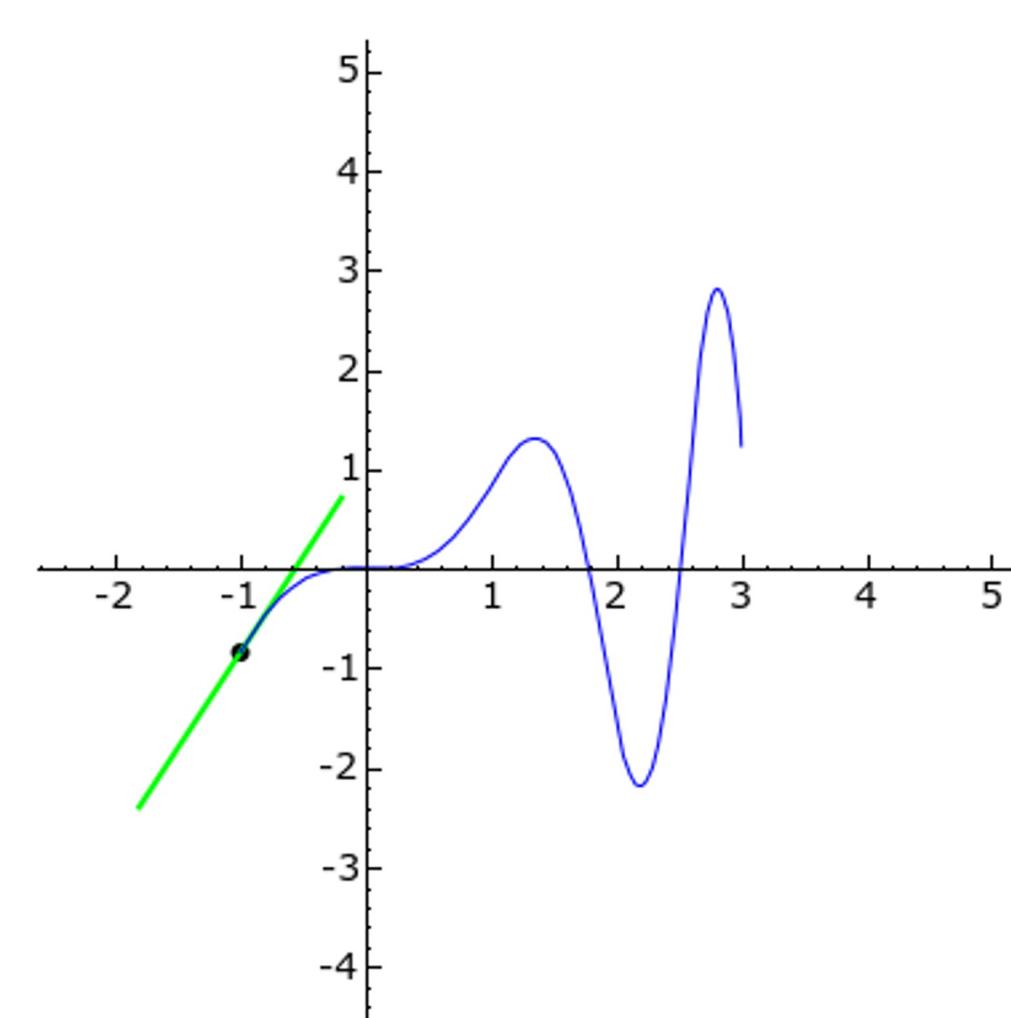
3. Calculate
predictions

4. Calculate
loss

5. Minimize

6. Go

7. Do this over and over for several epochs (one epoch is one pass through all the training data), until you meet some chosen stopping criteria



Define a neural network model with an architecture and a set of trainable parameters (weights and biases)

1. Init.

2. Col.

3. Cal.
pred.

4. Cal.
loss.

5. Min.

6. Go

7. Do this over and
over until the chosen stopping



01

Machine learning: **function approximation** based on **training**

02

We pick the model family \mathcal{F} . The parameters Θ are found automatically through training

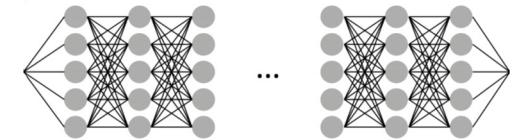
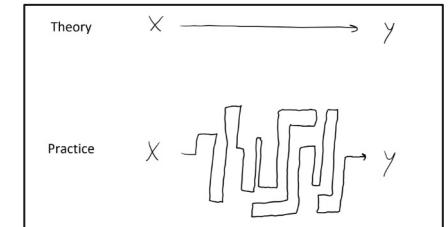
03

In practice: some models are better suited to some tasks than others.

04

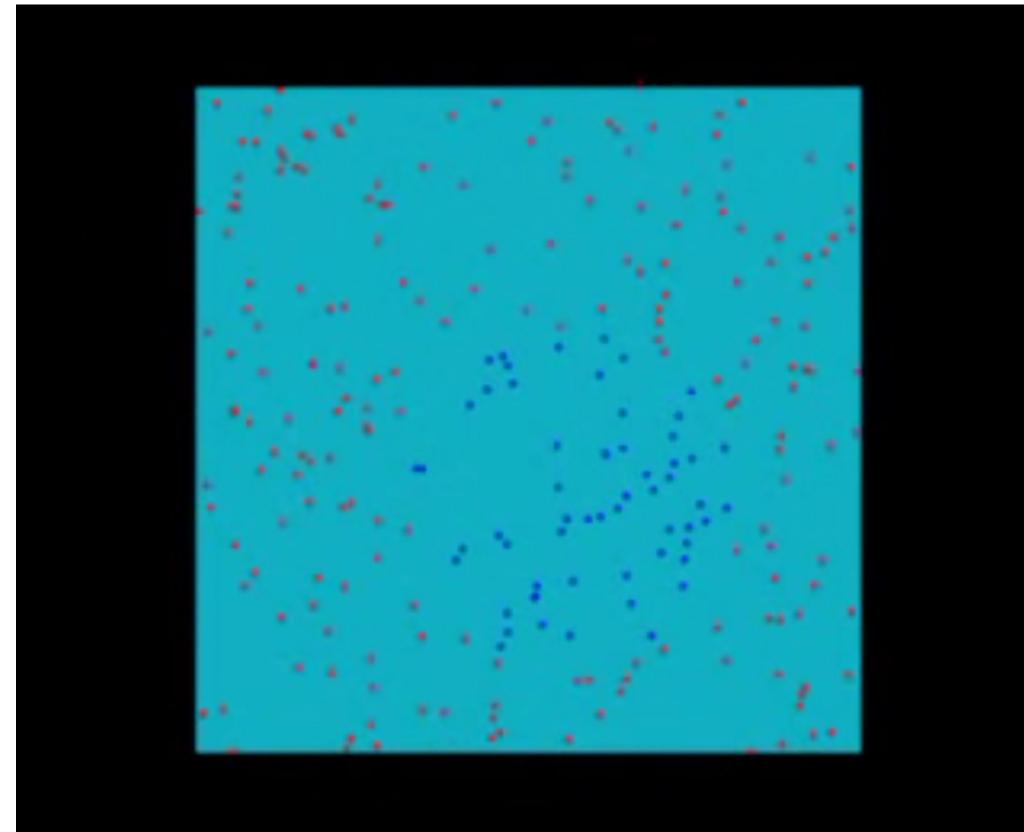
In deep learning: a particular choice of \mathcal{F} . A **composition** of simple functions, trained jointly, typically using GPUs or other accelerators.

$$y \approx f(x; \theta)$$

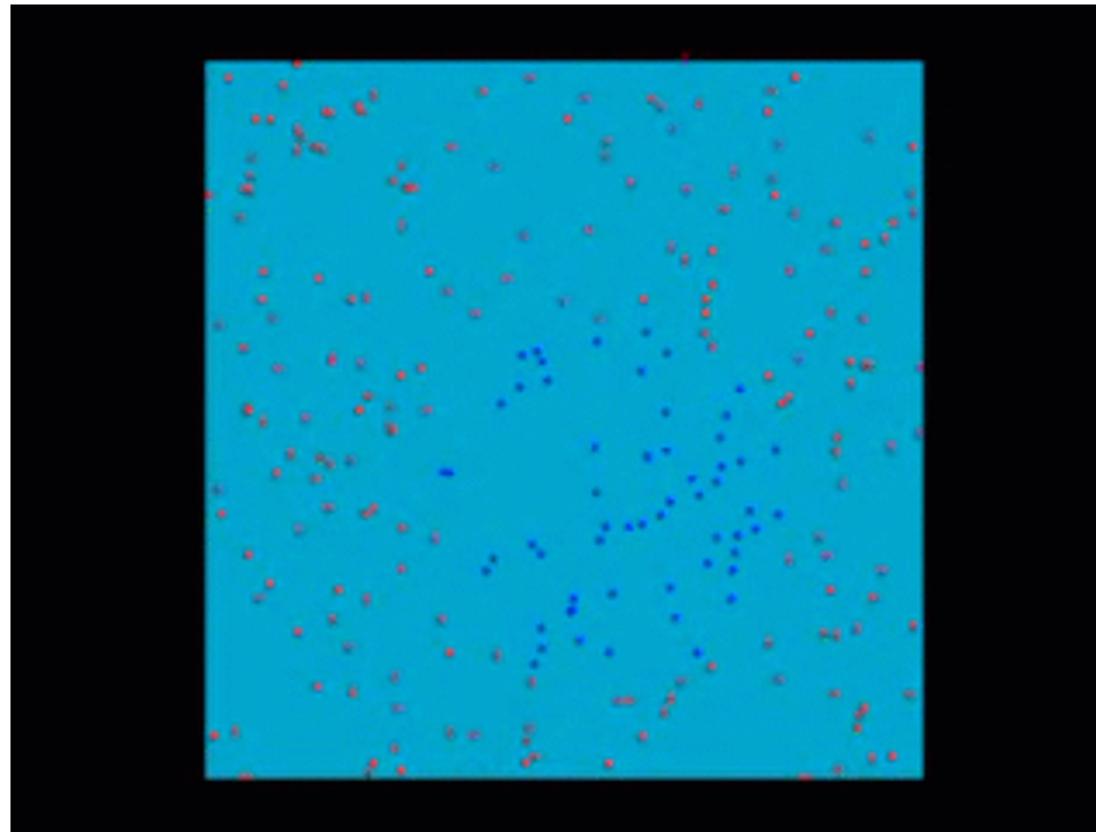


Representations and the power of transformations

Transformations

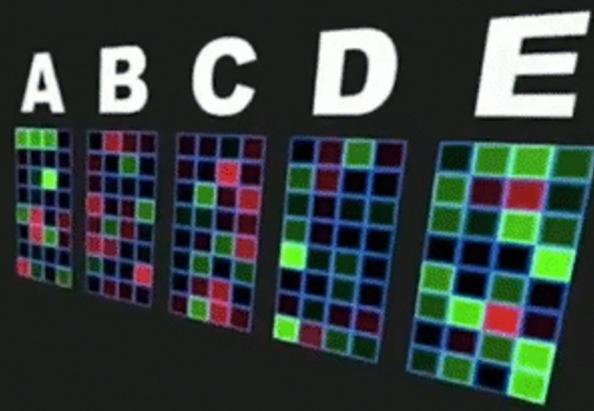


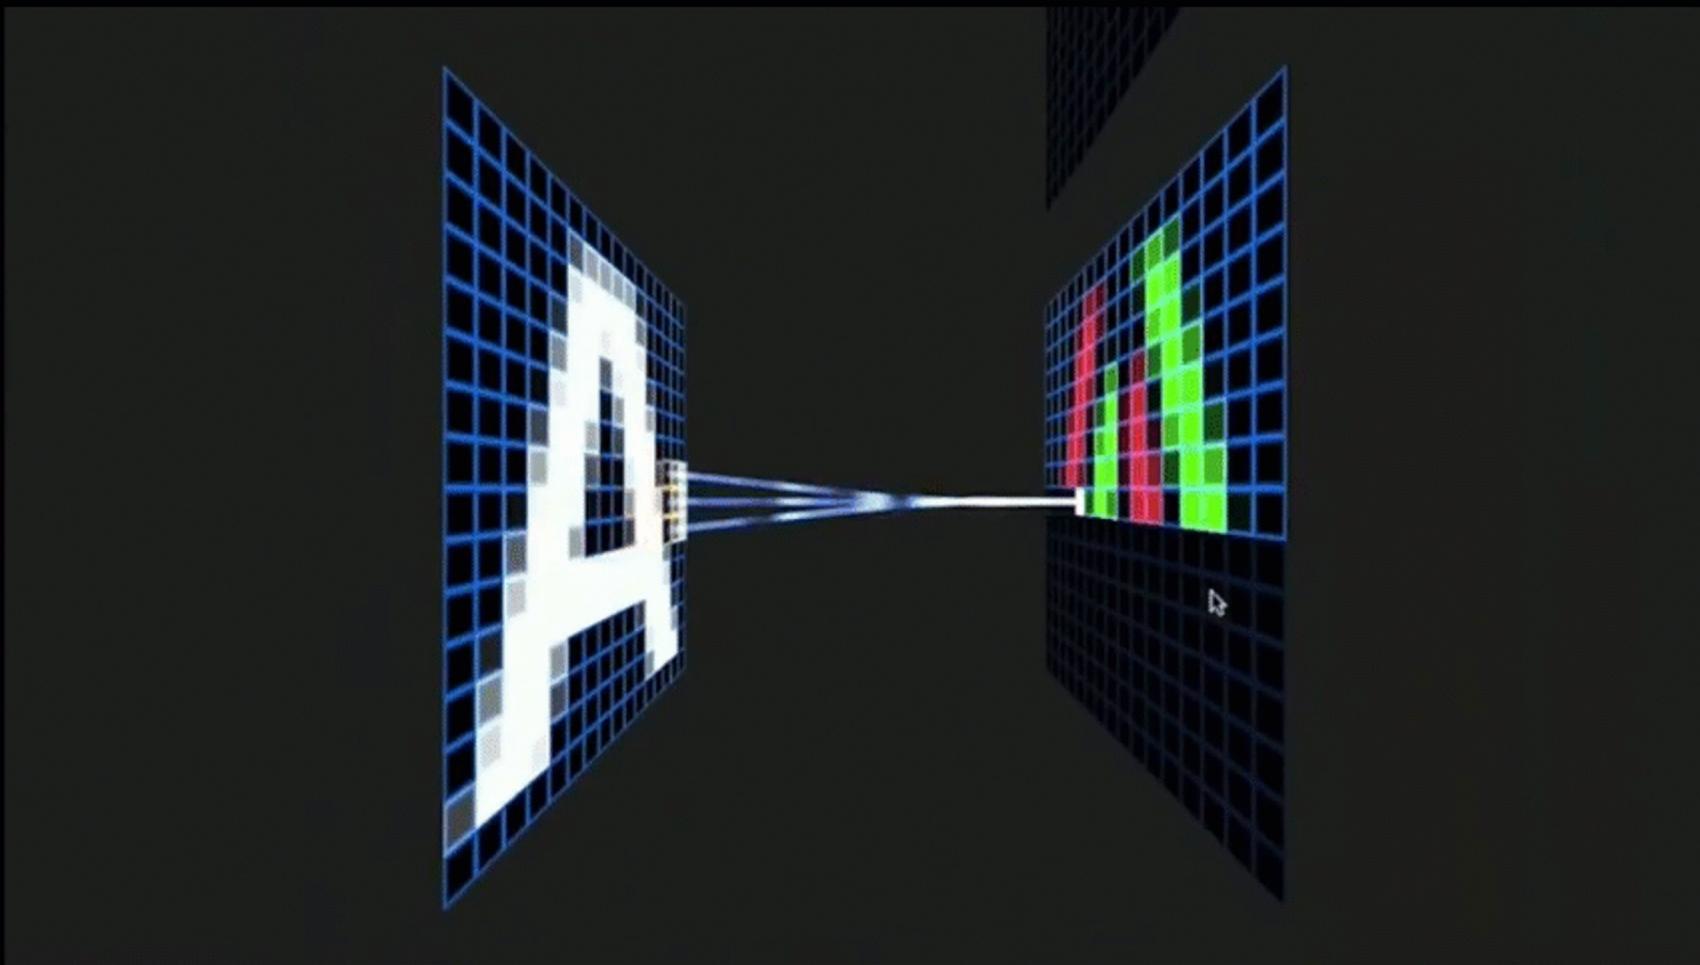
Transformation



What transformations should be applied?

In deep learning: ***transformations providing useful representations automatically found via training***





01

Machine learning: **function approximation** based on **training**

02

We pick the model family F . The parameters Θ are found automatically through training

03

In practice: some models are better suited to some tasks than others.

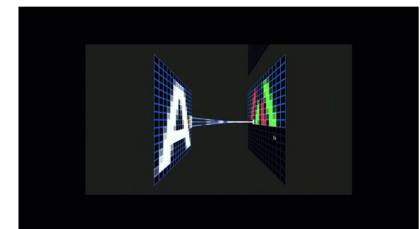
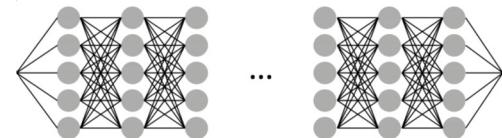
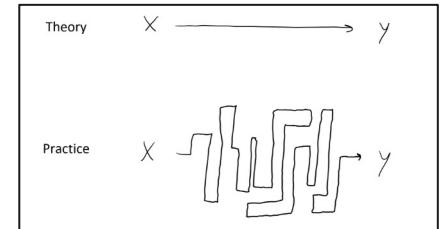
04

In deep learning: a particular choice of F . A **composition** of simple functions, trained jointly, typically using GPUs or other accelerators.

05

In deep learning: finding a **hierarchical set of geometric transformations**, ending in a representation that makes the desired task easy to solve.

$$y \approx f(x; \theta)$$



01

Machine learning: **function approximation** based on **training**

$$y \approx f(x; \theta)$$

02

We pick the model family \mathcal{F} . The parameters Θ are found automatically through training

03

In practice: some models are better suited to some tasks than others.

Deep learning

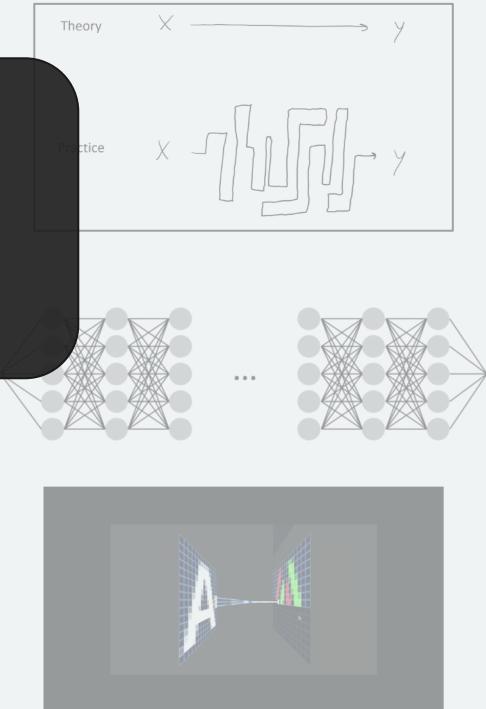
searching for good hierarchical geometric representations

04

In deep learning: a particular choice of \mathcal{F} . A **composition** of simple functions, trained jointly, typically using GPUs or other accelerators.

05

In deep learning: finding a **hierarchical set of geometric transformations**, ending in a representation that makes the desired task easy to solve.



Plan



What *is* deep learning?

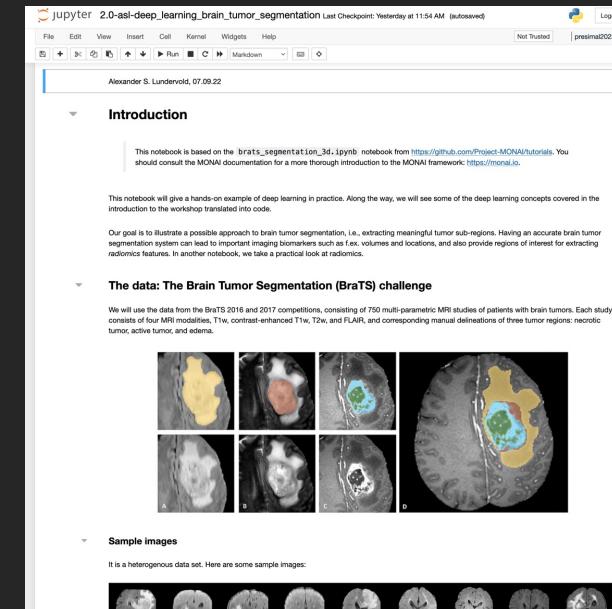
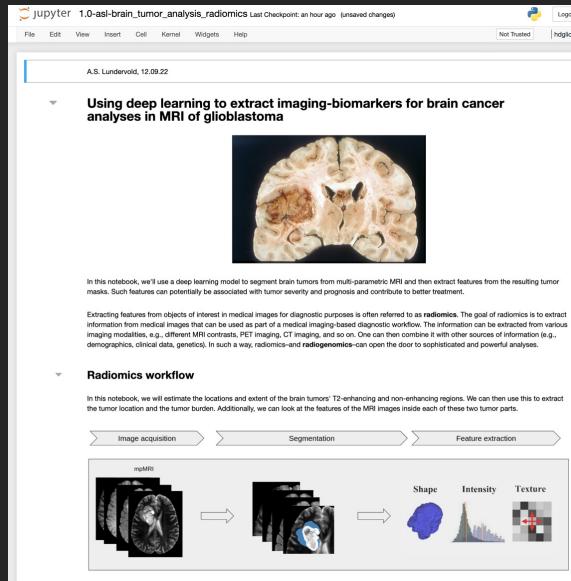


How do you *do* deep learning?



Some perspectives,
opportunities and
challenges

The workshop notebooks and fastMONAI provide a start. Please get in touch if you want additional pointers!



Installing